

TC3020 | Machine Learning (ML)

08 Artificial Neural Networks


Profr. Rafael Pérez Torres

r.perez@tec.mx

EiC | ITESM | Campus MTY


Non-linear hypothesis

- When dealing with non-linear decision boundaries it is possible to perform feature engineering to achieve separation of classes.
 - However, this is a risky operation, we might end up with way too many features.
 - Doing that operation by hand is cumbersome.
- **There should be a way to better select our features.**




Artificial Neural Networks (ANN)

- Origins: algorithms that try to mimic the human brain
- It is said that mammal brains follow a **single recipe to learn/process different types of stimuli**
- If that is the case, it could be possible to design an algorithm to learn or process any type of information




Model representation




Model representation

- **The perceptron**
- Parameters are called **weights**
- We still add a **bias unit**
- We need an activation function: sigmoid $\frac{1}{1+e^{-\theta^T x}}$
 - 0 output: inhibited neuron
 - 1 output: excited neuron


$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \theta_3 \end{bmatrix}$$



ANN architecture



ANN's activations



$a_i^{(j)}$ = activation of unit i in layer j

$\theta^{(j)}$ = matrix of weights controlling the function mapping from layer j to layer $j + 1$

θ_{23} : arrives to neuron 2, coming from neuron 3

$$\Theta^{(1)} \in \mathbb{R}^{3 \times 4}$$

It arrives to 3 neurons coming out from 4 neurons

$$a_1^{(2)} = g(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3)$$

$$a_2^{(2)} = g(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3)$$

$$a_3^{(2)} = g(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3)$$

$$h_{\Theta}(x) = a_1^{(3)} = g(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)})$$


If network has s_j units in layer j , s_{j+1} units in layer $j + 1$, then $\Theta^{(j)}$ will be of dimension $s_{j+1} \times (s_j + 1)$

E.g., for $\theta^{(2)}$. Layer 2 has 3 units, layer 3 has 1 unit. $\theta^{(2)}$ is of size 1×4

ANN's operation

- Working with a neural network involves two steps:
 - **Feed forward**: to predict $h_{\theta}(x)$
 - **Backpropagation**: to adjust the weights in θ :
 - This is the actual core of the training for an ANN.
 - When a good value of θ has been found, the model could be exported easily (just the matrix of weights θ plus the calculation of $h_{\theta}(x)$ is enough to predict the class for new examples).
- As we are dealing with matrixes, using vector operations can speed up and clarify the feed forward and backpropagation processes.

ANN vectorized implementation



$$a_1^{(2)} = g(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3)$$

$$a_2^{(2)} = g(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3)$$

$$a_3^{(2)} = g(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3)$$

$$h_{\Theta}(x) = g(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)})$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad z^{(2)} = \begin{bmatrix} z_1^{(2)} \\ z_2^{(2)} \\ z_3^{(2)} \end{bmatrix}$$

$$z^{(2)} = \Theta^{(1)} x$$

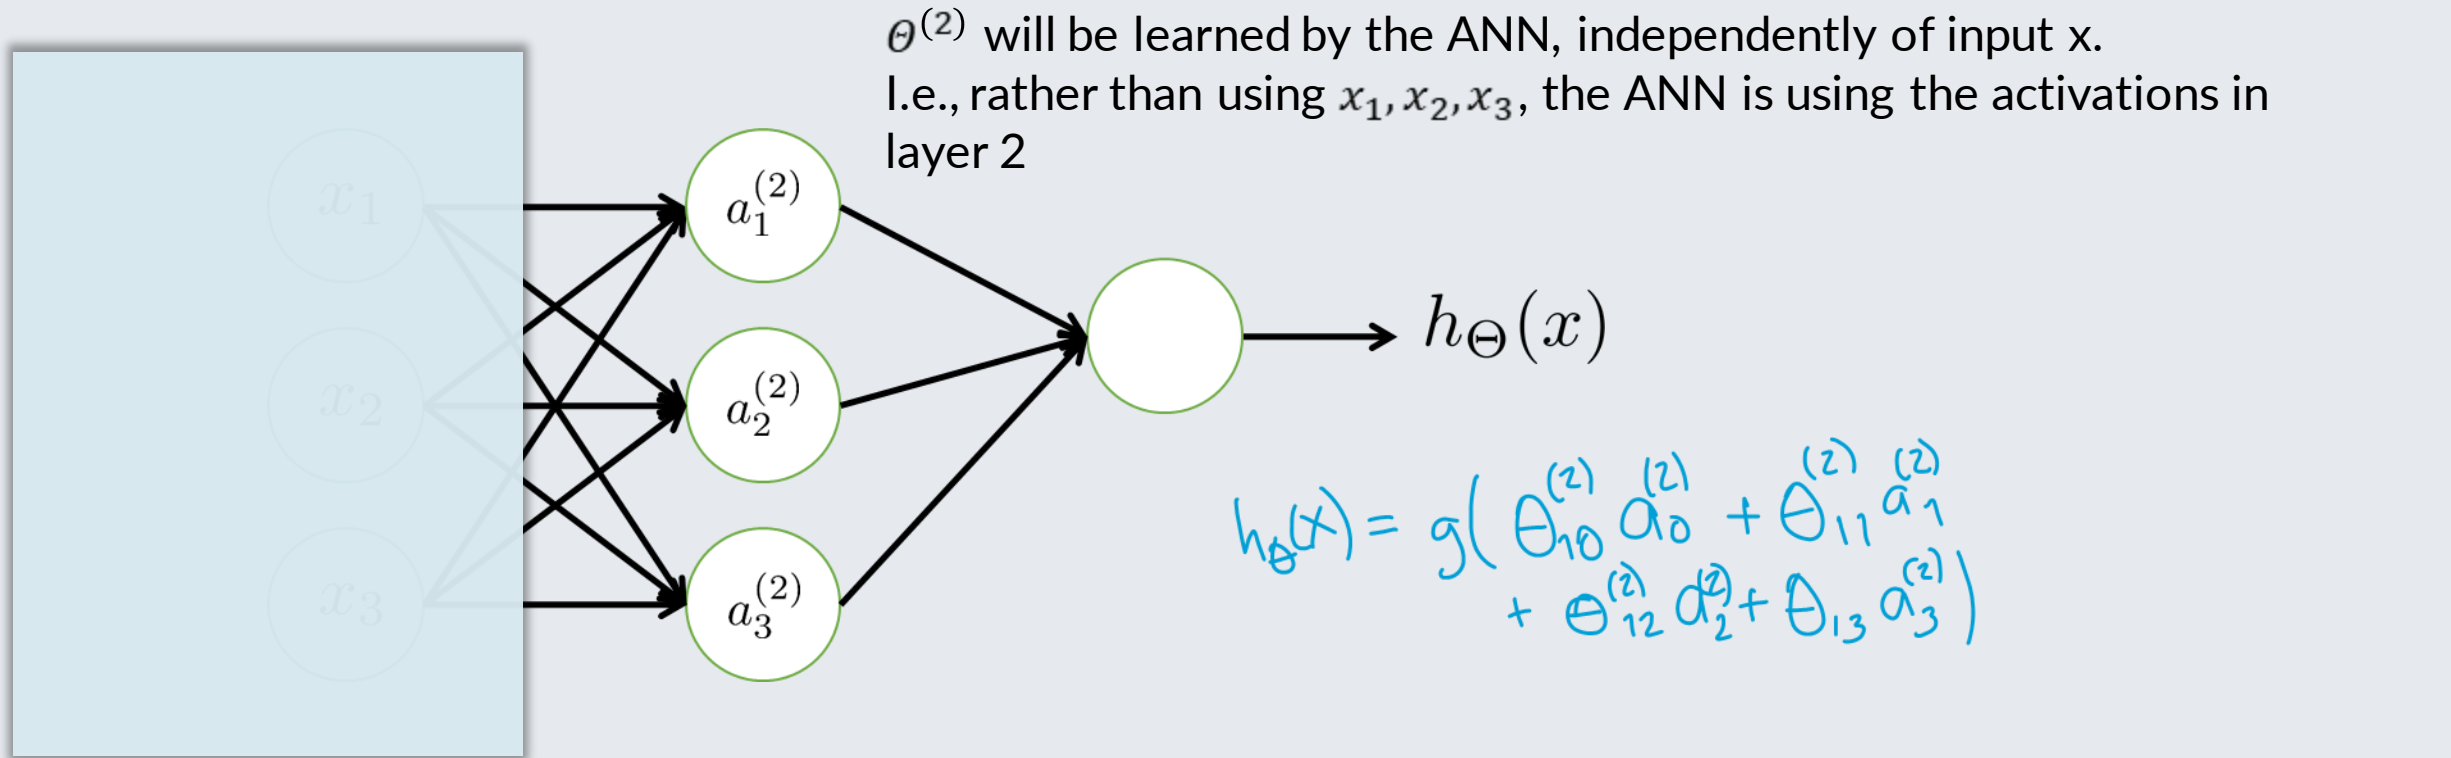
$$a^{(2)} = g(z^{(2)})$$

$$\text{Add } a_0^{(2)} = 1$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$h_{\Theta}(x) = a^{(3)} = g(z^{(3)})$$

ANN learning its own features





- **This is a lot like Logistic Regression!**
- Take each perceptron (neuron) individually, we find Logistic Regressions.
- With lots of them, it is clear how ANNs produce non-linear boundaries.
- Somehow the features become something more complex during the process.

Non-linear classification example

- “XNOR” function: negating the XOR


| x1 | x2 | Xor | Xnor |
|----|----|-----|------|
| 0 | 0 | 0 | 1 |
| 0 | 1 | 1 | 0 |
| 1 | 0 | 1 | 0 |
| 1 | 1 | 0 | 1 |




Example: Simple AND gate

$$x_1, x_2 \in \{0, 1\}$$

$$y = x_1 \text{ AND } x_2$$




$$h_{\theta}(x) = g(-30 + 20x_1 + 20x_2)$$



| x_1 | x_2 | $h_{\Theta}(x)$ |
|-------|-------|-----------------|
| 0 | 0 | |
| 0 | 1 | |
| 1 | 0 | |
| 1 | 1 | |

Example: Simple OR gate




$$h_{\theta}(x) = g(-10 + 20x_1 + 20x_2)$$

| x_1 | x_2 | $h_{\Theta}(x)$ |
|-------|-------|-----------------|
| 0 | 0 | |
| 0 | 1 | |
| 1 | 0 | |
| 1 | 1 | |

Example: Simple NOR gate

(NOT x_1) AND (NOT x_2)





1 iff both are = 0



$$h_{\theta}(x) = g(10 - 20x_1 - 20x_2)$$

| x_1 | x_2 | $h_{\Theta}(x)$ |
|-------|-------|-----------------|
| 0 | 0 | |
| 0 | 1 | |
| 1 | 0 | |
| 1 | 1 | |

Example: Putting it all together for XNOR



| x_1 | x_2 | $a_1^{(2)}$ | $a_2^{(2)}$ | $h_{\Theta}(x)$ |
|-------|-------|-------------|-------------|-----------------|
| 0 | 0 | 0 | 1 | 1 |
| 0 | 1 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 |
| 1 | 1 | 1 | 0 | 1 |

| x_1 | x_2 | Xor | Xnor |
|-------|-------|-----|------|
| 0 | 0 | 0 | 1 |
| 0 | 1 | 1 | 0 |
| 1 | 0 | 1 | 0 |
| 1 | 1 | 0 | 1 |

Multiclass classification



Pedestrian




Car



Motorcycle



Truck



$$h_{\Theta}(x) \in \mathbb{R}^4$$

$$h_{\Theta}(x) \approx \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

when pedestrian

$$h_{\Theta}(x) \approx \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

when
car

$$h_{\Theta}(x) \approx \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

when motorcycle

etc

Summary

- In this session we introduced Artificial Neural Networks.
 - Its inspiration in nature.
 - Its architecture.
 - The basic process to predict a class.
 - How it can create non-linear separations.
 - How it could be linked to Logistic Regression.
 - How it helps to follow a one-vs-all approach.